



DIGITAL SIGNAL & IMAGE MANAGEMENT


Giannelli Alessio Imbonati Lorenzo Valoti Davide

SUMMARY

TASK 1: Audio Signal Classification -
Speech Commands Dataset

TASK 2: Image Classification -
FGVC Aircraft 100 Dataset using Densenet201

TASK 3: Image Retrieval -
Feature Re-Weighting in CBIR



TASK 1 - Dataset & Pre-Processing



Info about dataset

Dataset size \approx 2GB

N. of instances = 64720

N. of instances per class \approx 2150

30 Class Labels

```
[array(['bed', 'bird', 'cat', 'dog', 'down', 'eight', 'five', 'four', 'go',  
      'happy', 'house', 'left', 'marvin', 'nine', 'no', 'off', 'on',  
      'one', 'right', 'seven', 'sheila', 'six', 'stop', 'three', 'tree',  
      'two', 'up', 'wow', 'yes', 'zero'], dtype='<U6'))]
```


Splitting

TRAIN (80%)

TEST (20%)

Features Extraction

```
def feats_mfcc(input, sample_rate):  
    cepstral_features = mfcc(input, sample_rate, numcep=20) #Features extraction  
    zero_vector = np.zeros((100-cepstral_features.shape[0], cepstral_features.shape[1]))  
    cepstral_features = np.vstack((cepstral_features, zero_vector))  
    return cepstral_features
```



Example of
MFCC
Feature

TASK 1 - Model Architecture



Layers and Parameters

```
net.summary()
```

```
Model: "model_2"
```

Layer (type)	Output Shape	Param #
input_3 (InputLayer)	[(None, 100, 20)]	0
gru_2 (GRU)	(None, 50)	10800
dense_2 (Dense)	(None, 30)	1530

```
=====
```

```
Total params: 12,330
```

```
Trainable params: 12,330
```

```
Non-trainable params: 0
```

Dropout = 0.3


Activation = Softmax

TASK 1 - Model Performance

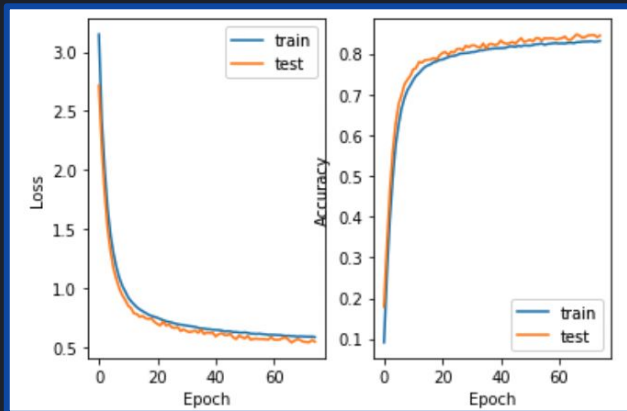
Classification Report

	precision	recall	f1-score	support
0.0	0.89	0.92	0.90	391
1.0	0.94	0.92	0.93	351
2.0	0.92	0.94	0.93	341
3.0	0.88	0.91	0.90	345
4.0	0.86	0.90	0.88	484
5.0	0.95	0.92	0.93	492
accuracy			0.91	12943
macro avg	0.92	0.91	0.91	12943
weighted avg	0.91	0.91	0.91	12943

Confusion Matrix



Loss & Accuracy Trend



Very high performance, touching 90% accuracy on both training and validation set.

TASK 1 - Model Evaluation



Audio Prediction

```
np.argmax(previsione)

10

encoder.inverse_transform([[10]])

array([[ 'house' ]], dtype='<U6')
```

Distribution Probability

```
previsione

array([[3.9753129e-05, 5.0238521e-07, 1.1391650e-06, 2.7455920e-03,
        4.6533391e-06, 3.7695609e-02, 1.2968192e-07, 2.0860985e-05,
        2.8214217e-04, 8.0605365e-05, 4.8249230e-01, 3.7027390e-03,
        2.0234948e-08, 7.9701730e-04, 1.9732230e-04, 3.9731449e-01,
        1.7278563e-02, 1.7386535e-02, 4.9062535e-05, 1.8000244e-05,
        3.3243868e-02, 7.7199639e-04, 3.0823336e-03, 2.3951443e-04,
        6.7771517e-04, 1.0602005e-04, 5.4005993e-04, 7.4405796e-07,
        1.1608218e-03, 6.9894508e-05]], dtype=float32)
```

The model correctly predicts audio with the "house" class with a probability of almost 50%.

TASK 2 - Dataset & Pre-Processing

Info about Dataset

Dataset Size: 2.76 GB

Number of classes: 100

Number of Instances per class: 100

Splitting

TRAIN (60%)
6000 instances

VALIDATION (30%)
3000 instances

TEST (10%)
1000 instances

Cropping



TASK 2 - Model Architecture V1

Transfer Architecture DenseNet201

```
base_net = keras.applications.DenseNet201(input_shape=(224,224,3),weights='imagenet', include_top=False)
for layer in base_net.layers[:501]:
    layer.trainable = False
```



Layer (type)	Output Shape	Param #
input_5 (InputLayer)	[(None, 224, 224, 3)]	0
tf.math.truediv_4 (TFOpLambda)	(None, 224, 224, 3)	0
tf.nn.bias_add_2 (TFOpLambda)	(None, 224, 224, 3)	0
tf.math.truediv_5 (TFOpLambda)	(None, 224, 224, 3)	0
densenet201 (Functional)	(None, 7, 7, 1920)	18321984
average_pooling2d_2 (AveragePooling2D)	(None, 3, 3, 1920)	0
global_average_pooling2d_2 (GlobalAveragePooling2D)	(None, 1920)	0
dense_6 (Dense)	(None, 1024)	1967104
batch_normalization_4 (Batch Normalization)	(None, 1024)	4096
dropout_4 (Dropout)	(None, 1024)	0
dense_7 (Dense)	(None, 512)	524800
batch_normalization_5 (Batch Normalization)	(None, 512)	2048
dropout_5 (Dropout)	(None, 512)	0
dense_8 (Dense)	(None, 100)	51300


Total params: 20,871,332
Trainable params: 9,055,396
Non-trainable params: 11,815,936

Transfer Learning

TASK 2 - Model ArchitectureV2

Transfer Architecture DenseNet201

```
base_net = keras.applications.DenseNet201(input_shape=(224,224,3),weights='imagenet', include_top=False)
base_net.trainable = True
```



Layer (type)	Output Shape	Param #
input_4 (InputLayer)	[(None, 224, 224, 3)]	0
tf.math.truediv_2 (TFOpLambda)	(None, 224, 224, 3)	0
tf.nn.bias_add_1 (TFOpLambda)	(None, 224, 224, 3)	0
tf.math.truediv_3 (TFOpLambda)	(None, 224, 224, 3)	0
densenet201 (Functional)	(None, 7, 7, 1920)	18321984
average_pooling2d_1 (AveragePooling2D)	(None, 3, 3, 1920)	0
global_average_pooling2d_1 (GlobalAveragePooling2D)	(None, 1920)	0
dense_3 (Dense)	(None, 1024)	1967104
batch_normalization_2 (Batch Normalization)	(None, 1024)	4096
dropout_2 (Dropout)	(None, 1024)	0
dense_4 (Dense)	(None, 512)	524800
batch_normalization_3 (Batch Normalization)	(None, 512)	2048
dropout_3 (Dropout)	(None, 512)	0
dense_5 (Dense)	(None, 100)	51300


Total params: 20,871,332
Trainable params: 20,639,204
Non-trainable params: 232,128

Transfer Learning

TASK 2 - Model ArchitectureV3

Transfer Architecture ResNet50

```
base_net = keras.applications.ResNet50(input_shape=(224,224,3),weights='imagenet', include_top=False)
base_net.trainable = True
```



Layer (type)	Output Shape	Param #
input_9 (InputLayer)	[(None, 224, 224, 3)]	0
tf.__operators__.getitem_1 (SlicingOpLambda)	(None, 224, 224, 3)	0
tf.nn.bias_add_4 (TFOpLambda)	(None, 224, 224, 3)	0
resnet50 (Functional)	(None, 7, 7, 2048)	23587712
average_pooling2d_4 (AveragePooling2D)	(None, 3, 3, 2048)	0
global_average_pooling2d_4 (GlobalAveragePooling2D)	(None, 2048)	0
dense_12 (Dense)	(None, 1024)	2098176
batch_normalization_8 (BatchNormalization)	(None, 1024)	4096
dropout_8 (Dropout)	(None, 1024)	0
dense_13 (Dense)	(None, 512)	524800
batch_normalization_9 (BatchNormalization)	(None, 512)	2048
dropout_9 (Dropout)	(None, 512)	0
dense_14 (Dense)	(None, 100)	51300

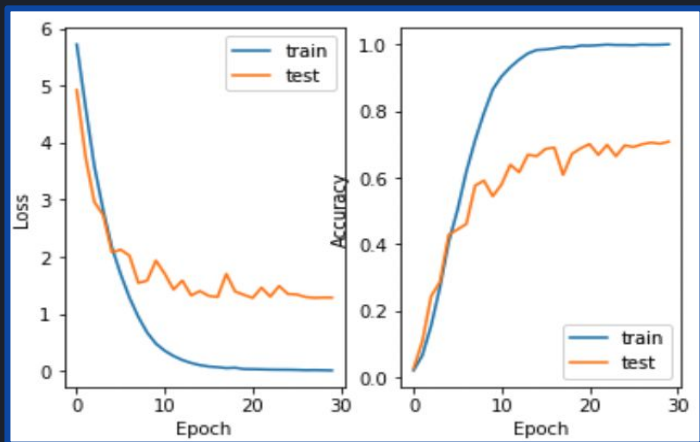
Total params: 26,268,132
Trainable params: 26,211,940
Non-trainable params: 56,192

Transfer Learning

TASK 2 - Model Performance V1

DenseNet201 9M trainable parameters

Loss & Accuracy Trend



The performance of the model is fair since it achieves almost 100% on the training set against 65-70% accuracy on the validation set.

It is emphasized, however, that from the 10th epoch onward there is a risk of overfitting problem.

Test loss	Test accuracy
1.289	70.75%

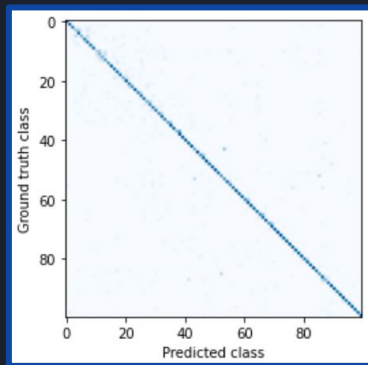
TASK 2 - Model Performance V2

DenseNet201 20M trainable parameters


Classification Report

	precision	recall	f1-score
0	0.86	0.97	0.91
1	0.93	0.85	0.89
2	0.96	0.74	0.83
19	0.65	0.52	0.58
20	0.73	0.94	0.82
21	0.65	0.73	0.69
22	0.82	0.70	0.75
23	0.97	0.91	0.94
accuracy			0.81
macro avg	0.82	0.81	0.80
weighted avg	0.82	0.81	0.80

Confusion Matrix



Loss & Accuracy Trend



The performance of the model is very good as it achieves 80% accuracy on the validation set.

The trend shows some performance fluctuations but in general it is quite stable and the validation follows the growth of the training set.

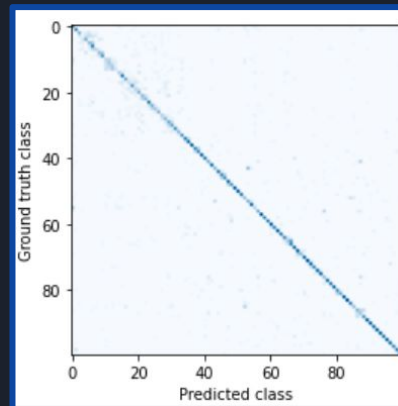
TASK 2 - Model Performance V3

ResNet50 26M trainable parameters

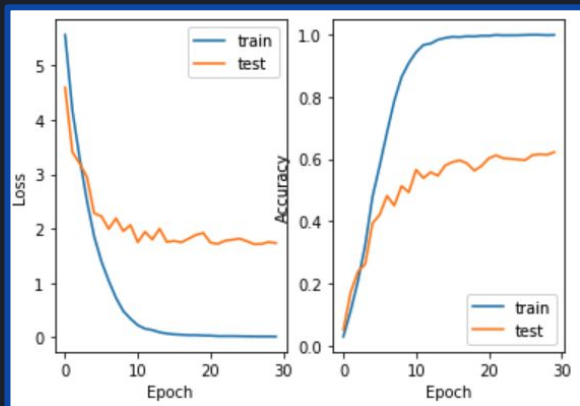
Classification Report

	precision	recall	f1-score
0	0.34	0.67	0.45
1	0.67	0.67	0.67
2	0.56	0.44	0.49
97	0.92	0.70	0.79
98	0.74	0.74	0.74
99	0.82	0.82	0.82
accuracy			0.61
macro avg	0.62	0.61	0.61
weighted avg	0.62	0.61	0.61

Confusion Matrix



Loss & Accuracy Trend



The performance of the model is not acceptable since it achieves almost 100% accuracy on the training set but settles around 60% on the validation set. In general, it seems that the model suffers from overfitting.

TASK 2 - Model Evaluation

on Web Image

Distribution Probability

```
previsione
```

```
array([[8.45307895e-08, 8.59294147e-09, 2.28450489e-07, 1.04331457e-05,  
        5.10848849e-07, 9.47635385e-07, 9.56761141e-08, 1.16615745e-04,  
        1.19382069e-01, 6.69421115e-06, 1.71462332e-07, 1.86927124e-07,  
        3.83096435e-08, 1.61001561e-04, 8.33525717e-01, 1.18540612e-03,
```

on Test Set

Test loss	Test accuracy
0.8971	80.30%

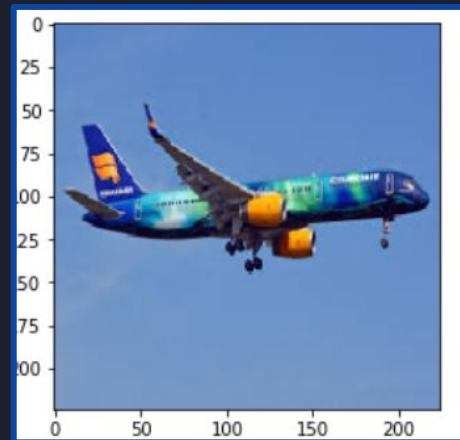
Image Prevision

```
np.argmax(previsione)
```

```
14
```

```
train.class_names[14]
```

```
'757-200'
```





TASK 3 - Feature Re-Weighting in CBIR

Implementation of the following paper

Feature Re-weighting in Content-Based Image Retrieval

Gita Das¹, Sid Ray¹, and Campbell Wilson²

¹ Clayton School of Information Technology
Monash University
Victoria 3800, Australia

{Gita.Das, Sid.Ray}@csse.monash.edu.au

² Caulfield School of Information Technology
Monash University
Victoria 3800, Australia

Campbell.Wilson@csse.monash.edu.au

Main concepts:

- Use of the previous neural network as feature extractor
- Use of weighted Minkowski distance as similarity measure
- Update of the query results according to user preferences

TASK 3 - Feature Extraction


Load Task2 Model

```
temp = keras.models.load_model('Model/densenet201_final_task2.h5')  
  
layer_name = 'dense_1'  
newmodel = Model(inputs=temp.input, outputs=temp.get_layer(layer_name).output)  
newmodel.summary()
```

Splitting

TRAIN
6000 instances


TEST
1000 instances



Features normalization

$$f'_i = \frac{f_{i,org} - \mu_i}{3\sigma_i}$$

$$f_i = \frac{f'_i + 1}{2}$$



TASK 3 - Image Retrieval (Query)

Manhattan similarity measure

$$D(I, Q) = \sum_{i=1}^M w_i * |f_{iI} - f_{iQ}|$$

Weights are constant for the first round of retrieval

Top20 Accuracy on test set
77,56%

Images Similarity to the test image



TASK 3 - Rebalancing type 1

Update weights formula Type 1

$$\text{weight - type1 : } w_i^{k+1} = \frac{\epsilon + \sigma_{N_r, i}^k}{\epsilon + \sigma_{rel, i}^k}, \epsilon = 0.0001$$

New weight for the i-th feature is equal to the division between the standard deviation over the 20 retrieved images and the standard deviation over the relevant images at the previous round

$$\mathbf{w}^{k+1} = 0.9 * \mathbf{w}^k + 0.1 * \mathbf{w}^{k+1}$$

Round number	Top20 Precision
Round 0	77.56
Round 1	83.94
Round 2	84.56
Round 3	85.10
Round 4	85.41
Round 5	85.54

TASK 3 - Rebalancing type 2

Update weights formula Type 2

$$\text{weight} - \text{type2} : w_i^{k+1} = \frac{\delta_i^k}{\epsilon + \sigma_{rel,i}^k}$$

$$\delta_i^k = 1 - \frac{\sum_{l=1}^k |\psi_i^{l,U}|}{\sum_{l=1}^k |F_i^{l,U}|}$$

New weight for the i-th feature is equal to the division between the sigma quantity defined in the second formula, that depends on the **dominant range**, and the standard deviation over the relevant images at the previous round

$$\mathbf{w}^{k+1} = 0.9 * \mathbf{w}^k + 0.1 * \mathbf{w}^{k+1}$$

Round number	Top20 Precision
Round 0	77.56
Round 1	61.70
Round 2	58.84
Round 3	59.91
Round 4	60.09
Round 5	60.53



TASK 3 - Rebalancing type 3

Update weights formula Type 3


$$\text{weight - type3 : } w_i^{k+1} = \delta_i^k * \frac{\epsilon + \sigma_{N_r,i}^k}{\epsilon + \sigma_{rel,i}^k}$$

New weight for the i-th feature is equal to the the delta value defined in the previous slide by the weights of type 1

$$\mathbf{w}^{k+1} = 0.9 * \mathbf{w}^k + 0.1 * \mathbf{w}^{k+1}$$

Round number	Top20 Precision
Round 0	77.56
Round 1	60.33
Round 2	57.35
Round 3	57.94
Round 4	57.85
Round 5	57.77

TASK 3 - Rebalancing Types Trend



Type 1 rebalancing is definitely the best since it shows increasing growth.


The other two types of rebalancing do not produce any improvement.

Let's leave room for the demo...


Aircraft image search engine

Nessun file selezionato


Query:




Results:




737-200




737-200




737-200



737-200



737-200



737-200